

# Mnist\_with\_keras

June 15, 2025

```
[2]: from keras.datasets import mnist
```

```
(X_train, y_train), (X_test, y_test) = mnist.load_data()
```

Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz>

11490434/11490434 1s

0us/step

```
[3]: X_train.shape , y_train.shape , X_test.shape , y_test.shape
```

```
[3]: ((60000, 28, 28), (60000,), (10000, 28, 28), (10000,))
```

```
[4]: X_train.min(), X_train.max(), X_test.min(), X_test.max()
```

```
[4]: (np.uint8(0), np.uint8(255), np.uint8(0), np.uint8(255))
```

```
[5]: import matplotlib.pyplot as plt
```

```
plt.figure(figsize=(5 ,5))
```

```
plt.imshow(X_train[0], cmap='gray')
```

```
[5]: <matplotlib.image.AxesImage at 0x7c523f9cf210>
```



0.	, 0.	, 0.	, 0.	, 0.	,
0.	, 0.	, 0.	, 0.	, 0.	,
0.	, 0.	, 0.	, 0.	, 0.	,
0.	, 0.	, 0.	, 0.	, 0.	,
0.	, 0.	, 0.	, 0.	, 0.	,
0.	, 0.	, 0.	, 0.	, 0.	,
0.	, 0.	, 0.	, 0.	, 0.	,
0.	, 0.	, 0.	, 0.	, 0.	,
0.	, 0.	, 0.	, 0.	, 0.	,
0.	, 0.	, 0.	, 0.	, 0.	,
0.	, 0.	, 0.	, 0.	, 0.	,
0.	, 0.	, 0.	, 0.	, 0.	,
0.	, 0.	, 0.	, 0.	, 0.	,
0.	, 0.	, 0.	, 0.	, 0.	,
0.	, 0.	, 0.	, 0.	, 0.	,
0.	, 0.	, 0.	, 0.	, 0.	,
0.	, 0.	, 0.	, 0.	, 0.	,
0.	, 0.	, 0.	, 0.	, 0.	,
0.	, 0.	, 0.	, 0.	, 0.	,
0.	, 0.	, 0.	, 0.	, 0.	,
0.	, 0.	, 0.	, 0.	, 0.	,
0.	, 0.	, 0.	, 0.	, 0.	,
0.	, 0.	, 0.	, 0.	, 0.	,
0.	, 0.	, 0.01176471,	0.07058824,	0.07058824,	
0.07058824,	0.49411765,	0.53333333,	0.68627451,	0.10196078,	
0.65098039,	1.	, 0.96862745,	0.49803922,	0.	,
0.	, 0.	, 0.	, 0.	, 0.	,
0.	, 0.	, 0.	, 0.	, 0.	,
0.	, 0.11764706,	0.14117647,	0.36862745,	0.60392157,	
0.66666667,	0.99215686,	0.99215686,	0.99215686,	0.99215686,	
0.99215686,	0.88235294,	0.6745098	, 0.99215686,	0.94901961,	
0.76470588,	0.25098039,	0.	, 0.	, 0.	,
0.	, 0.	, 0.	, 0.	, 0.	,
0.	, 0.	, 0.	, 0.19215686,	0.93333333,	
0.99215686,	0.99215686,	0.99215686,	0.99215686,	0.99215686,	
0.99215686,	0.99215686,	0.99215686,	0.98431373,	0.36470588,	
0.32156863,	0.32156863,	0.21960784,	0.15294118,	0.	,
0.	, 0.	, 0.	, 0.	, 0.	,
0.	, 0.	, 0.	, 0.	, 0.	,
0.	, 0.07058824,	0.85882353,	0.99215686,	0.99215686,	
0.99215686,	0.99215686,	0.99215686,	0.77647059,	0.71372549,	
0.96862745,	0.94509804,	0.	, 0.	, 0.	,
0.	, 0.	, 0.	, 0.	, 0.	,
0.	, 0.	, 0.	, 0.	, 0.	,
0.	, 0.	, 0.	, 0.	, 0.	,
0.31372549,	0.61176471,	0.41960784,	0.99215686,	0.99215686,	
0.80392157,	0.04313725,	0.	, 0.16862745,	0.60392157,	
0.	, 0.	, 0.	, 0.	, 0.	,

0. , 0. , 0. , 0. , 0. ,  
 0. , 0. , 0. , 0. , 0. ,  
 0. , 0. , 0. , 0. , 0.05490196 ,  
 0.00392157, 0.60392157, 0.99215686, 0.35294118, 0. ,  
 0. , 0. , 0. , 0. , 0. ,  
 0. , 0. , 0. , 0. , 0. ,  
 0. , 0. , 0. , 0. , 0. ,  
 0. , 0. , 0. , 0. , 0. ,  
 0. , 0. , 0. , 0. , 0.54509804 ,  
 0.99215686, 0.74509804, 0.00784314, 0. , 0. ,  
 0. , 0. , 0. , 0. , 0. ,  
 0. , 0. , 0. , 0. , 0. ,  
 0. , 0. , 0. , 0. , 0. ,  
 0. , 0. , 0. , 0. , 0. ,  
 0. , 0. , 0.04313725, 0.74509804, 0.99215686 ,  
 0.2745098 , 0. , 0. , 0. , 0. ,  
 0. , 0. , 0. , 0. , 0. ,  
 0. , 0. , 0. , 0. , 0. ,  
 0. , 0. , 0. , 0. , 0. ,  
 0. , 0. , 0. , 0. , 0. ,  
 0. , 0.1372549 , 0.94509804, 0.88235294, 0.62745098 ,  
 0.42352941, 0.00392157, 0. , 0. , 0. ,  
 0. , 0. , 0. , 0. , 0. ,  
 0. , 0. , 0. , 0. , 0. ,  
 0. , 0. , 0. , 0. , 0. ,  
 0. , 0. , 0. , 0. , 0. ,  
 0.31764706, 0.94117647, 0.99215686, 0.99215686, 0.46666667 ,  
 0.09803922, 0. , 0. , 0. , 0. ,  
 0. , 0. , 0. , 0. , 0. ,  
 0. , 0. , 0. , 0. , 0. ,  
 0. , 0. , 0. , 0. , 0. ,  
 0. , 0. , 0. , 0. , 0.17647059 ,  
 0.72941176, 0.99215686, 0.99215686, 0.58823529, 0.10588235 ,  
 0. , 0. , 0. , 0. , 0. ,  
 0. , 0. , 0. , 0. , 0. ,  
 0. , 0. , 0. , 0. , 0. ,  
 0. , 0. , 0. , 0. , 0. ,  
 0. , 0. , 0. , 0.0627451 , 0.36470588 ,  
 0.98823529, 0.99215686, 0.73333333, 0. , 0. ,  
 0. , 0. , 0. , 0. , 0. ,  
 0. , 0. , 0. , 0. , 0. ,  
 0. , 0. , 0. , 0. , 0. ,  
 0. , 0. , 0. , 0. , 0. ,  
 0. , 0. , 0. , 0.97647059, 0.99215686 ,  
 0.97647059, 0.25098039, 0. , 0. , 0. ,  
 0. , 0. , 0. , 0. , 0. ,  
 0. , 0. , 0. , 0. , 0. ,

5

```
[10]: y_train[0]
```

```
[11]: from tensorflow.keras.models import Sequential
      from tensorflow.keras.layers import Dense
      from tensorflow.keras.utils import to_categorical
```

```
[13]: y_train.shape , y_test.shape
```

```
[14]: y_train[0]
```

```
Model: "sequential"
```

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 64)	50,240
dense_1 (Dense)	(None, 64)	4,160
dense_2 (Dense)	(None, 10)	650

Total params: 55,050 (215.04 KB)

Trainable params: 55,050 (215.04 KB)

Non-trainable params: 0 (0.00 B)

## 1 Model Summary

First Hidden Layer \* No. of inputs: 784 \* No. of neurons in hidden layer 1: 64 \* Total No. of Parameters:  $784 * 64 + 64 = 50240$

Second Hidden Layer \* No. of inputs: 64 \* No. of neurons in hidden layer 2: 64 \* Total No. of Parameters:  $64 * 64 + 64 = 4160$

Output Layer \* No. of inputs: 64 \* No. of neurons in output layer: 10 \* Total No. of Parameters:  $64 * 10 + 10 = 650$

Total No. of Parameters:  $50240 + 4160 + 650 = 55050$

## 2 Alternative Model Building

1. Define then add layers with activation

```
model = Sequential()
model.add(Dense(64, activation='rule', input_shape=(784,)))
model.add(Dense(64, activation='rule'))
model.add(Dense(10, activation='softmax'))
```

2. Define then add layers without activation

```
from tensorflow.keras.layers import Activation
model = Sequential()
model.add(Dense(64, input_dim=784))
model.add(Activation('rule'))
model.add(Dense(64))
model.add(Activation('rule'))
```

```
model.add(Dense(10))
model.add(Activation('softmax'))
```

### 3. Functional API syntax

```
from tensorflow.keras.models import Model
from tensorflow.keras.layers import Input, Dense
inputs = Input(shape=(784,))
x = Dense(64, activation='relu')(inputs)
x = Dense(64, activation='relu')(x)
predictions = Dense(10, activation='softmax')(x)
model = Model(inputs=inputs, outputs=predictions)
```

```
[17]: # Train the model
model.fit(X_train, y_train, epochs=10, batch_size=32, validation_split=0.3)
print('Training Completed \n')

# Evaluate the model
loss, accuracy = model.evaluate(X_test, y_test)
print(f'Test Loss: {loss:.4f}')
print(f'Test Accuracy: {accuracy * 100:.2f}%')
```

```
Epoch 1/10
1313/1313          6s 4ms/step -
accuracy: 0.9936 - loss: 0.0194 - val_accuracy: 0.9776 - val_loss: 0.0918
Epoch 2/10
1313/1313          11s 4ms/step -
accuracy: 0.9943 - loss: 0.0173 - val_accuracy: 0.9744 - val_loss: 0.1065
Epoch 3/10
1313/1313          11s 5ms/step -
accuracy: 0.9958 - loss: 0.0144 - val_accuracy: 0.9779 - val_loss: 0.0940
Epoch 4/10
1313/1313          10s 5ms/step -
accuracy: 0.9960 - loss: 0.0138 - val_accuracy: 0.9754 - val_loss: 0.1053
Epoch 5/10
1313/1313          9s 4ms/step -
accuracy: 0.9947 - loss: 0.0148 - val_accuracy: 0.9720 - val_loss: 0.1307
Epoch 6/10
1313/1313          6s 5ms/step -
accuracy: 0.9937 - loss: 0.0170 - val_accuracy: 0.9752 - val_loss: 0.1172
Epoch 7/10
1313/1313          5s 4ms/step -
accuracy: 0.9971 - loss: 0.0092 - val_accuracy: 0.9746 - val_loss: 0.1245
Epoch 8/10
1313/1313          10s 4ms/step -
accuracy: 0.9972 - loss: 0.0091 - val_accuracy: 0.9738 - val_loss: 0.1309
Epoch 9/10
1313/1313          5s 4ms/step -
accuracy: 0.9953 - loss: 0.0133 - val_accuracy: 0.9742 - val_loss: 0.1289
```



Epoch 10/10  
1313/1313 9s 4ms/step -  
accuracy: 0.9972 - loss: 0.0090 - val\_accuracy: 0.9755 - val\_loss: 0.1324  
Training Completed

313/313 1s 2ms/step -  
accuracy: 0.9691 - loss: 0.1569  
Test Loss: 0.1310  
Test Accuracy: 97.24%

```
[18]: # Batch vs Epoch
batch_size = 32
training_samples = 60000
# Calculate the number of batches
batches_per_epoch = training_samples / batch_size
batches_per_epoch
```

[18]: 1875.0

### 3 Some Karas Options

```
[16]: # Activation function
from tensorflow.keras.activation import relu, sigmoid, linear, softmax, tanh

# Optimizer
from tensorflow.keras.optimizers import SGD, RMSprop, Adam, Adadelta, Adagrad,
↳Adamax, Nadam, Ftrl

# loss function
from tensorflow.keras.losses import mean_squared_error, mean_absolute_error,
↳binary_crossentropy, categorical_crossentropy,
↳sparse_categorical_crossentropy, kullback_leibler_divergence, poisson,
↳cosine_similarity

# metrics
from tensorflow.keras.metrics import accuracy, binary_accuracy,
↳categorical_accuracy, sparse_categorical_accuracy
```

```
[21]: def build_model(hidden_activation="relu",
↳output_activation='softmax', optimizer='Adam',
↳loss='categorical_crossentropy'):
    model = Sequential()
    model.add(Dense(64, input_shape=(784,), activation=hidden_activation))
    model.add(Dense(64, activation=hidden_activation))
    model.add(Dense(10, activation=output_activation))
    model.compile(optimizer=optimizer, loss=loss, metrics=['accuracy'])
    return model
```

```
[22]: my_model = build_model()
      my_model.summary()
```

Model: "sequential\_2"

Layer (type)	Output Shape	Param #
dense_6 (Dense)	(None, 64)	50,240
dense_7 (Dense)	(None, 64)	4,160
dense_8 (Dense)	(None, 10)	650

Total params: 55,050 (215.04 KB)

Trainable params: 55,050 (215.04 KB)

Non-trainable params: 0 (0.00 B)

```
[29]: history = my_model.fit(X_train, y_train, epochs=10, batch_size=32,
                             ↪validation_split=0.3)
```

```
Epoch 1/10
1313/1313          7s 5ms/step -
accuracy: 0.9915 - loss: 0.0254 - val_accuracy: 0.9713 - val_loss: 0.1153
Epoch 2/10
1313/1313          9s 7ms/step -
accuracy: 0.9940 - loss: 0.0185 - val_accuracy: 0.9704 - val_loss: 0.1322
Epoch 3/10
1313/1313          5s 4ms/step -
accuracy: 0.9948 - loss: 0.0166 - val_accuracy: 0.9707 - val_loss: 0.1333
Epoch 4/10
1313/1313          6s 5ms/step -
accuracy: 0.9949 - loss: 0.0154 - val_accuracy: 0.9696 - val_loss: 0.1427
Epoch 5/10
1313/1313          5s 4ms/step -
accuracy: 0.9949 - loss: 0.0150 - val_accuracy: 0.9721 - val_loss: 0.1402
Epoch 6/10
1313/1313          6s 5ms/step -
accuracy: 0.9964 - loss: 0.0106 - val_accuracy: 0.9673 - val_loss: 0.1652
Epoch 7/10
1313/1313          9s 4ms/step -
accuracy: 0.9952 - loss: 0.0145 - val_accuracy: 0.9701 - val_loss: 0.1519
```

```
Epoch 8/10
1313/1313          6s 5ms/step -
accuracy: 0.9968 - loss: 0.0096 - val_accuracy: 0.9694 - val_loss: 0.1652
Epoch 9/10
1313/1313          11s 5ms/step -
accuracy: 0.9975 - loss: 0.0078 - val_accuracy: 0.9688 - val_loss: 0.1706
Epoch 10/10
1313/1313          10s 5ms/step -
accuracy: 0.9938 - loss: 0.0179 - val_accuracy: 0.9682 - val_loss: 0.1824
```

```
[25]: batch_size = 32
      training_samples = 0.8 * 60000
      iterations_per_epoch = training_samples / batch_size
      iterations_per_epoch
```

```
[25]: 1500.0
```

```
[30]: history.params
```

```
[30]: {'verbose': 'auto', 'epochs': 10, 'steps': 1313}
```

```
[31]: history.history.keys()
```

```
[31]: dict_keys(['accuracy', 'loss', 'val_accuracy', 'val_loss'])
```


```
[32]: import pandas as pd
      history_df = pd.DataFrame(history.history)
      history_df
```

```
[32]:
```

	accuracy	loss	val_accuracy	val_loss
0	0.991595	0.025370	0.971333	0.115310
1	0.992929	0.020782	0.970444	0.132237
2	0.993381	0.020074	0.970722	0.133330
3	0.994190	0.017717	0.969611	0.142737
4	0.994476	0.016430	0.972056	0.140246
5	0.995333	0.014707	0.967333	0.165164
6	0.994976	0.014446	0.970111	0.151878
7	0.995095	0.014481	0.969389	0.165226
8	0.997762	0.007059	0.968833	0.170586
9	0.994524	0.015622	0.968167	0.182351


```
[33]: # plot the loss
      import seaborn as sns
      sns.set_style()
      sns.lineplot(data=history_df[['loss', 'val_loss']])
```

```
[33]: <Axes: >
```



```
[34]: # plot the accuracy
sns.lineplot(data=history_df[['accuracy', 'val_accuracy']])
```

[34]: <Axes: >



[ ]: